Bringing together visual analytics and probabilistic programming languages

Jonas Aaron Gütter

Friedrich Schiller Universität Jena

Matrikelnr 152127

Prof.Dr. Joachim Giesen

M. Sc. Phillip Lucas

26. November 2018

Zusammenfassung

A probabilistic programming language (PPL) provides methods to represent a probabilistic model by using the full power of a general purpose programming language. Thereby it is possible to specify complex models with a comparatively low amount of code. With Uber, Microsoft and DARPA focusing research efforts towards this area, PPLs are likely to play an important role in science and industry in the near future. However in most cases, models built by PPLs lack appropriate ways to be properly visualized, although visualization is an important first step in detecting errors and assessing the overall fitness of a model. This could be resolved by the software Lumen, developed by Philipp Lucas, which provides several visualization methods for statistical models. PPLs are not yet supported by Lumen, and the goal of the master thesis at hand is to change that by implementing an interface between Lumen and a chosen PPL, so that exploring PPL models by visual analytics becomes possible. The thesis will be divided into two main parts, the first part being an overview about how PPLs work and what existing PPLs there are available. Out of these, the most appropriate one will be chosen for the task. The second, more practical part will then document the actual implementation of the interface.

Inhaltsverzeichnis

1 Road Map

1. Getting started

   - set up Master thesis document
   - Probabilistic Programming Languages
     - play at least with: PyMC3, Stan
     - read the docs, wiki, ...
     - download the libraries
     - reproduce the getting started tutorials
     - –> understand the ideas and how to use it, get a feeling for it
   - theoretic background: Read Bayesian Data Analysis part I, Chapter 1,2 and part II, chapter 6,7
   - Lumen
     - install locally and play with
     - understand main idea of Lumen and what we want to do with it
   - Start filling up your MA thesis document
     - understand and write down in MA thesis the "why & what"
     - describe the background of the work, e.g. summarize PPLs
   - give a short presentation
     - what have you done and learned
     - what do you plan to do?
     - why is it relevant?
     - how do you plan to measure the success?

2. First connection of PPLs to Lumen

   - Start from small, very simple and specific example. Generalize later.
   - Choose PPL to work with
     - work out requirements on PPL
     - work out preferred features of PPL
     - choose a PPL based on these requrements and preferences
   - design wrapper of PPL with Lumen
     - work out requirement and interface
     - identify necessary work on Lumen
     - identify necessary work
   - Connect chosen specific example with lumen
   - Continue to work on your master thesis document!

3. Improve, generalize and clean up the connection of your PPL to Lumen

2   Introduction

In this section, the motivation for the thesis is explained, the task of the thesis is made clear. Furthermore, the structure of the thesis is outlined here. It is possible to get the basic ideas and results of the thesis from reading only this section and the conclusion.

3   Concepts of Bayesian Statistics

Theoretical concepts of building statistical models using a Bayesian approach are explained in this section. Also graphical model checking is treated here, as that will be the common use case for the Lumen interface created in this thesis.

3.1   Rules of Probabilistic Inference

When working with Bayesian models, it will be necessary to transform conditional, marginal and joint distributions into one another. There are three rules of probabilistic inference which achieve this: The chain rule, the total probability rule, and the Bayes' rule. The following explanations are taken from [2].

3.1.1   Chain rule

The chain rule is used to calculate a joint probability distribution of several variables from local conditional probability distributions of these variables:

$$P(X_1, X_2, ...X_n) = P(X_1)P(X_2|X_1)P(X_3|X_1, X_2)...P(X_n|X_1, X_2, ...X_{n-1})) \tag{1}$$

3.1.2   Total probability rule

The total probability rule calculates the probability distribution over a subset of variables, also called a marginal distribution, by summing out all the other variables, that is by summing the probability distributions for each combination of values of these variables:

$$P(\boldsymbol{X}|\boldsymbol{Z}) = \sum_{\boldsymbol{y}} P(\boldsymbol{X}, \boldsymbol{Y} = \boldsymbol{y}|\boldsymbol{Z}) \tag{2}$$

3.1.3   Bayes' rule

Bayes' rule calculates the probability of a cause, given an effect, by using the prior probability of the cause and the probability of the effect, given the cause.

$$P(X|Y) = (P(Y|X) * P(X))/P(Y) \tag{3}$$

3.2   General principle

Goal of Bayesian Statistics: Setting up a probability model over model parameters and observed data

When doing Bayesian data analysis, we have observed data Y and a prior belief $P(\theta)$ about the mechanism that generated the data, and we use those to set up a full probability model $P(y, \theta|Y)$ over all model parameters θ and all possible data y. The conceptual difference to conventional Statistics is two-fold: First, we don't use prior beliefs in conventional Statistics (at least not for parameter estimation), second, in conventional Statistics we only set up a probability model over the data, not over the model parameters. The model parameters in conventional Statistics are fixed after the parameter estimation.

The principle of setting up the full probability model of Bayesian Statistics (also called the joint probability distribution) is as follows:
According to the chain rule, we can compute the full probability model with the following formula:
$$p(y, \theta|Y) = p(y|\theta, Y) * p(\theta|Y) \tag{4}$$

Since the probability for a specific value of y does not depend on the observed data Y, if we already know θ, we can write equivalently:

$$p(y, \theta|Y) = p(y|\theta) * p(\theta|Y) \tag{5}$$

$p(y|\theta)$ is usually easy to compute since it is given by the chosen model class. For computing the second term, $P(\theta|Y)$, we use the Bayes rule:

$$p(\theta|Y) = p(Y|\theta) * p(\theta)/p(Y) \tag{6}$$

This is where the prior belief comes into play. $p(\theta)$ represents the prior belief, the distribution over the model parameters θ, that we assume from prior knowledge. $p(Y|\theta)$ is the likelihood: The probability that the observed data occur at a given parameter. The problems of choosing a suitable likelihood and prior are treated in sections 3.3 and 3.4. $p(Y)$, the overall probability of the observed data, does not have to be computed since it can be treated as a constant. (TODO: UNDERSTAND THIS BETTER, HOW p(Y) CAN BE IGNORED)

3.3   Choosing an appropriate prior distribution

There are two interpretations of prior distributions. The *population* interpretation, where the prior distribution is thought of as a population of possible parameters, from where the current parameter is drawn. This, as far as I understand, requires the range of possible values to be known, e.g from past experience. On the other hand, the *state of knowledge* interpretation looks at the prior distribution as an expression of the user's knowledge or uncertainty, so that the assumption is plausible, that the real parameter value is taken from a random realization of the prior distribution.

When we have lots of knowledge about the parameter already, it makes sense to choose an *i*nformative prior, which has a big influence on the posterior distribution.On the other hand, if one does not want prior beliefs to affect the outcome of an analysis, one should not choose informative priors, even when lots of knowledge is available. E.g. when testing a hypothesis, it is not wise to include information in the prior that supports the hypothesis for fairness reasons(?? look that up again on p. 56.). If there is no sufficient data to estimate a prior, it is desirable to choose a prior that is *n*oninformative, meaning that it will contribute very little to the posterior distribution ('let the data speak for itself'). A common noninformative prior, motivated by Laplace's principle of insufficient reason, is the uniform distribution. However, the uniform distribution is not ideal since it is dependent on the parametrization: Applying a uniform distribution to $p(x)$ leads to a non-uniform distribution when looking at $p(x^2)$ and the other way round [3]. Jeffrey's approach to find noninformative prior distributions negates this problem by choosing a prior that is invariant to the parameterization (see https://eventuallyalmosteverywhere.wordpress.com/2013/05/10/bayesian-inference-and-the-jeffreys-prior/). Besides informative and noninformative prior distributions there is also the *weakly informative* prior distribution. This kind of prior does affect the posterior distribution in terms of regularization (e.g. it prevents extreme outliers), but it does not contain any further special knowledge about the parameter. A normal distribution with high variance is often used as weakly informative prior.

The parameters of the prior distributions are called hyperparameters. The property that prior and posterior distribution are of the parametric form (e.g., both are a beta distribution) (for a given likelihood distribution), is called conjugacy. Conjugate prior distributions have the advantages of being computationally convenient and being intertpretable as additional data.

A prior distribution is called *proper* if it does not depend on data and sums to 1 [3]. Proper distributions can be normalized. The uniform prior for example is *improper*, since it can't be normalized.

3.4   Sampling distribution

The sampling distribution, also called the likelihood of the data.
Often it is about which distribution class should be chosen for the prior and the likelihood. Standard, convenient distributions for single-parameter models: normal, binomial, Poisson, exponential. Those can also be combined to represent more complex distributions. For different classes of sampling distributions there are corresponding conjugate prior distributions which lead in turn to posterior distributions of the same form. [3]
distributions can be chosen for mathematical convenience. One could estimate hyperparameters of the likelihood distribution from the data in some cases. This is a bit of a circular reasoning, but apparently it is appropriate for [3].

3.5   Evaluating models

Since "all models are wrong", it is necessary, after a Bayesian model has been found, to capture the extent to which the model fails to describe reality. In Bayesian analysis, both the prior and the likelihood distribution can be sources of error.

Different models can lead to different posterior inferences, even if both models have a good fit.

model checking: How good does the model fit the data? Sensitivitiy analysis: Comparison of different models with good fit

This is necessary after a joint probability density and a posterior density are calculated. How good are the prior and the likelihood model? *Sensitivity analysis* deals with the question, how much the posterior is changed when one chooses a different (but also reasonable) model for the likelihood or the prior, or other assumptions affecting the posterior.

In theory: Setting up a 'super-model' as joint-distribution which includes all possible realities. Posterior of this super-model then automatically incorporates the sensitivity analysis. In practice this is conceptually and computationally infeasible and also is still dependent on the assumptions made when creating the super-model being correct.

One way is just to compare model predictions with the actual outcomes (this is referred to as external validation). If for some reason it is not possible to get the acutal outcomes (for example if they lie in the future), one needs to approximate external validation with the available data.

Posterior predictive checking: Use global summaries of the data distribution and the predictive distribution to evaluate the model: Draw random samples from the predictive distribution, compare e.g. the smallest of these values with the smallest value in the data...

model checking by calculating *test quantities*. *test quantities* are scalar summaries of parameters that are used to compare data to simulations. *test quantities* are similar to test statistics in the classical approach. Besides on the data, *test quantities* can also depend on the parameters.

Graphical comparison of simulated and data histograms (!! maybe here point out usefulness of Lumen)

tail-area-probability: p-value. p-value in classical statistics: Probability, that, given a parameter θ, a test statistic for replicated data is greater or equal the test statistic for the observed data,as shown in equation 7:

$$p_C = Pr(T(y^{rep}) >= T(y)|\theta) \tag{7}$$

[3]

The Bayesian p-value, on the other hand, is not conditional on a fixed θ, since in Bayesian statistics parameters are drawn from a distribution the same way as outcome variables. Instead, the observed data y is the fixed quantity for the Bayesian p-value, as shown in equation 8

$$p_B = Pr(T(y^{rep}, \theta) >= T(y, \theta)|y) \tag{8}$$

p-value depends on the chosen test quantity.

Bayesian predictive checking: choose and compute test quantites for the data and for simulations, then compute the p-value for the test quantity (probability that the test quantity of the data could have arisen under the model)

Test quantities give information about a specific aspect of the observed/simulated data. A model can predict values that are in some aspects similar to the data, in other aspects different. Therefore it can make sense to compute p-values fur several different test quantities. If a test quantity is dependent on the model parameters it has to be simulated both for the observed as well as for the predicted data. This makes it possible to compare both test statistics (for observed and for simulated values) pointwise, eg. in a scatterplot or in a histogram of the differences. The scatterplot should be symmetric around the 45 degree line, the histogram should include 0 (and be symmetric around 0??). Test quantities are most useful when they measure an aspect that is not directly recognizable in the probability model. Graphical visualization can be used to identify useful test quantities. If a test quantity has a p-value close to 0 or 1, the aspect that is analysed by the test quantity is not well captured in the model. Model checking depends on the data-collection model (how the data was collected), since data has to be replicated for computing a p-value. Posterior Inference however does not depend on the data collection model, as long as the likelihood model stays the same. Because of this, the same posterior results(what are the parameters) can get different model checking results (how well does the model fit) when the data-collection model is changed.

Marginal predictive checks: Instead of replicating data from the joint posterior distribution, one can also compute the probability for a certain outcome, and compare this probability with the data....(p.153).

Graphical posterior predictive checks: Basic Idea: Visualize the data alongside simulated data and so find discrepancies –> perfect in Lumen Possibilities:

- Direct display of all the data

- Display of data summaries (e.g. if the data is so large that it is not helpful to display all the data)

- Display a measure of discrepancy betweem model and data

MAYBE BUILD EXAMPLES WITH LUMEN TO ILLUSTRATE USEFULNESS OF LUMEN

Direct data display   Show the complete data and its correspondent simulations. It's best to show several replications.

Display data summaries

Display discrepancy measures   Bayesian residuals are defined by $y_i - g(x, \theta)$, where, in contrast to classical residuals, θ is not a point estimate of parameters, but a single draw from the posterior distribution of the parameters.

Difficult with discrete data, since residuals can only take certain discrete values. Graphical visualization can then be misleading since the discrete values can suggest patterns. To avoid this, one can use *binned residuals*, that is plotting the averages over a number of residuals and their corresponding outcomes.

4 Probabilistic Programming

In this section, the idea of Probabilistic Programming is explained. It is defined what a Probabilistic Programming Language is, and a number of PPLs are introduced. A focus is laid on the abilities of the PPLs for graphical model checking

"Probabilistic programming is an emerging paradigm in statistical learning, of which Bayesian modeling is an important sub-discipline."[1]
stochastic data types ($->$ probability distributions) make it easier to perform bayesian data analysis

4.1 What are Probabilistic Programming Languages

Modelle spezifizieren/beschreiben
Random Variables can be represented as objects
Posterior probability for parameters comes from the Bayes formula: $P(\theta|data) = p(data|\theta) * p(\theta)/p(data)$. The marginalized probability of the data $p(data)$ is given as follows: $p(data) = \int p(\theta, data)d\theta$. This integral can't be calculated (TODO: UNDERSTAND THAT BETTER), so PPLs instead sample from this distribution and get their inference from the samples.
[5]
effizienter in der Beschreibung von Modellen als herkömmliche Programmiersprachen [6]
unifying general purpose programming with probabilistic modeling [7]
A probabilistic reasoning system uses prior knowledge in the form of a probabilistic model to answer a certain query. The particular propertes of the query as well as the prior knowledge are given to an inference algorithm which returns an answer in the form of probabilities. Example is shown in figure 1. Probabilistic Programming is the implementation of a probabilistic reasoning system by using a programming language.

Traditional means for representing models are not always sufficient for probabilistic models. Therefore, probabilistic programming languages were introduced to be able to represent models with the full power of a programming language (http://www.probabilistic-programming.org/wiki/Home).

4.2 Inference through sampling

Using the Bayes formula, we get the posterior distribution the following way:

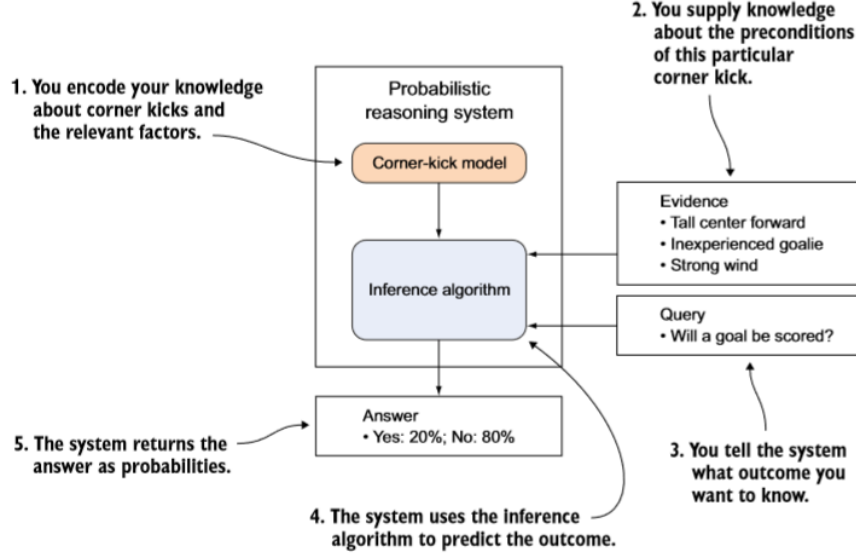$$p(\theta|x) = p(x|\theta) * p(\theta)/p(x) \tag{9}$$

Abbildung 1: General workflow example of a probabilistic reasoning system

We get both the likelihood $p(x|\theta)$ and the prior $p(\theta)$ from the model assumptions. The marginalized probability over the data, however, is more complicated. To compute it analytically, we would have to compute the following integral:

$$
\begin{aligned}
p(x) &= \int p(\theta, x)d\theta \\
&= \int p(x|\theta) * p(\theta)d\theta
\end{aligned}
\tag{10}
$$

Calculating this integral can be very difficult or impossible, since the integral can be multidimensional. So, it is not possible to get the posterior distribution of the parameters analytically. However, there are methods to sample exactly from the posterior distribution of the parameters. The easiest of them is called Simple Monte Carlo. There we approximate the integral by sampling over the likelihood and the prior and then taking the mean over the samples:

$$
\int p(x|\theta) * p(\theta)d\theta \approx 1/m * \sum_{}^{m} p(x|\theta_i), \quad \theta_i \sim p(\theta)
\tag{11}
$$

Simple Monte Carlo has a number of disadvantages, so more sophistically methods are genrally used, like Gibbs sampling or NUTS.

Markov Chain simulation:

Parameter values are drawn sequentially from an approximate distribution. In each sequence, the values depend only on the values of the preceding sequence

(and maybe on the index of the sequence), which makes those draws a markov chain.

So in each step, values are drawn from a transition distribution $T_t(\theta^t|\theta^{t-1})$.

The chain has the property of becoming more similar to the real posterior distribution of the parameters in each iteration.

4.3 Difference to conventional Programming Languages

When estimating parameters, not only the most likely value for the parameters is given, but also their uncertainty.

4.4 Comparing Different Probabilistic Programming Languages

- stan for python: https://pystan.readthedocs.io/en/latest/

- pymc3: https://docs.pymc.io/notebooks/getting_started.html#Case-study-2:-Coal-mining-disasters

- edward: http://edwardlib.org/getting-started

- pyro: http://pyro.ai/

4.4.1 Stan for python

Stan is an open-source program written in C++ that is designed for Bayesian inference on user-specified models. It consists of the following components:

- A modeling language that allows users to specify custom models

- An inference engine which uses Hamilton Monte Carlo methods for sampling to get an approximation for the posterior distribution

- An L-BFGS optimizer for finding local optima

- Procedures for automatic differentiation which are able to compute gradients required by the sampler and the optimizer

- Routines to monitor the convergence of parallel chains and compute inferences and effective sample sizes

- Wrappers for Python, R, Julia and other languages

The code for describing a model by using the aforementioned modeling language is divided into six blocks: *data*, *transformed data*, *parameters*, *transformed parameters*, *model*, and *generated quantities*: In *data* and *transformed data*, the structure of the input data, along with any constraints, is given. The difference between the former and the latter is that variables in *transformed data* are functions of other data variables, whereas variables in *data* are not. For example if in *data*, a variable x is listed, then in *transformed data* one can specify a variable

$y = x^2$. In *parameters* and *transformed parameters*, the parameters of the models are described, whereat, as in the data blocks, the latter contains variables that are functions of other parameters. In *model*, the model structure in the form of prior and likelihood distributions is given. Finally, the block *generated quantities* can be used to perform simulations and make predictions. In **??**, an example from [8] is shown, where the model $y = a_1 e^{-b_1 x} + a_2 e^{-b_2 x}$ is fit to data, using the PyStan interface of Stan for Python. Since not all of the aforementioned blocks are mandatory, only four of them are used in the example.
PyStan also provides basic plotting methods for the posterior distributions of the parameters.

is not able to perform inference on discrete parameters. Discrete data and discrete-data models, however, are possible

computes the log-posterior density

[8]

[9]

für den Code evtl. https://github.com/stephen-hoover/presentations zitieren

4.4.2   Pymc3

PyMC3 is an open-source probabilistic programming framework for Python. The following explanations are taken from [1]. Specification of Bayesian models in PyMC3 is done by encoding the prior, the sampling and the posterior distributions through three types of random variables: Stochastic, deterministic and observed stochastic ones. Stochastic random variables have values which are in part determined randomly, according to a chosen distribution. Commonly used probability distributions like Normal, Binomial etc. are available for this. Deterministic random variables, on the other hand, are not drawn from a distribution, but are calculated by fixed rules from other variables, for example by taking the sum of two variables. Lastly, there are the observed stochastic random variables which are similar to stochastic random variables, except that they get passed observed data as an argument, that should not be changed by any fitting algorithm. This kind of random variable can be used to represent sampling distributions.
PyMC3 mainly uses simulation techniques to draw inference on posterior distributions. It focuses especially on the No-U-Turn Sampler, a Markov Chain Monte Carlo algorithm, that relies on automated differentiation to get gradient information about continuous posterior distributions. PyMC3 also provides basic methods for plotting posterior distributions.
The code piece in **??** shows a simple example of a Bayesian model, taken from [1]. There, the data X1, X2 and Y is used to fit a regression model. First, prior distributions for the model parameters are set up as stochastic random variables, then the regression model itself is specified by a deterministic random variable and lastly the sampling distribution is described by an observed stochastic random variable to which the observed outcome Y is given as a parameter. Finally, the posterior distribution is simulated by drawing 500 samples from it.

```python
# Specify model
example_code = """
data {
// Define input data in this block
int N;
vector[N] x;
vector[N] y;
}
parameters {
// These are random parameters which we want to estimate
vector[2] log_a;
ordered[2] log_b;
real<lower=0> sigma;
}
transformed parameters {
// Create quantities derived from the parameters.
vector<lower=0>[2] a;
vector<lower=0>[2] b;
a <- exp(log_a);
b <- exp(log_b);
}
model {
// Define your model here
vector[N] ypred;
ypred <- a[1]*exp(-b[1]*x) + a[2]*exp(-b[2]*x);
y ~ lognormal(log(ypred), sigma);
log_a ~ normal(0,1);
log_b ~ normal(0,1);
}
"""

# Pass data to the model. x and y are the observed data
example_dat = {'x':x,'y':y,'N':len(x)}
# Fit model
sm = pystan.StanModel(model_code=example_code)
fit = sm.sampling(data=example_dat, iter=1000, chains=4)
print(fit)
```

Abbildung 2: Example code of a simple Bayesian model using Stan

```python
import pymc3 as pm

basic_model = pm.Model()

with basic_model:
# describe prior distributions of model parameters. Stochastic variables
alpha = pm.Normal('alpha', mu=0, sd=10)
beta = pm.Normal('beta', mu=0, sd=10, shape=2)
sigma = pm.HalfNormal('sigma', sd=1)
# specify model for the output parameter. Deterministic variable
mu = alpha + beta[0]*X1 + beta[1]*X2
# likelihood of the observations. Observed stochastic variable
Y_obs = pm.Normal('Y_obs', mu=mu, sd=sigma, observed=Y)

# model fitting by using sampling strategies
with basic_model:
# draw 500 posterior samples
trace = pm.sample(500)
pm.summary(trace)
```

Abbildung 3: Example code of a simple Bayesian model using PyMC3

strictly positive priors are transformed with log transformation, so that they are unconstrained, since that is better for sampling

There is the possibility to create own theano functions in python. Gradient based sampling methods don't work for user-defined functions however, except when a gradient is explicitly added.

Similarily pmc3 allows to define own distributions

4.4.3 Edward

4.4.4 Pyro

4.4.5 BUGS and Jags

mentioned in [8] based on graphical models

5 Lumen

In this section, the Software Lumen is introduced and its abilities and uses are elucidated, especially its potential uses for Bayesian model checking in conjunction with PPLs. This covers also the theoretical aspect: Which distribution do we want to visualize with Lumen and how do we get it (using the formula $P(y, \theta|Y) = P(y|\theta) * P(\theta|Y)$)? It is outlined, what the goal of the thesis is precisely, in other words what possibilities should the interface ideally give to a user who has written a model with a PPL. Too achieve this is goal it is necessary to

understand how Lumen works, so the functionality of Lumen is also explained here.

The overall goal of the thesis is to build an API, that makes it possible to a PyMC3 (or other PPL) user to get his model visualized in lumen with minimal effort. For example it would be good if one just has to call one function with a PyMC3 model as parameter to transport the model into the lumen frontend.

uses modelbase as backend: modelbase is a python package that provides means for fitting models to data and manipulate models by marginalization and conditionalization.

model classes are in `modelbase/mb_modelbase/models_core` data is in `mb_data`. model fitting is done in the script fit_models.py, which is also in `mb_data`. The corresponding model class has to be imported, the models have to be 'registered' in the known_models variable and then the models to be fit have to be listed in the variable debugincl.

Fitted models that can be accessed by lumen are then stored in `data_models`.

5.1 Functionality

In the modelbase repository, a model class is described by a python file that includes a class with several prescribed methods `https://ci.inf-i2.uni-jena.de/gemod/modelbase/blob/master/mb_modelbase/models_core/models.py` is a template for such a model class file, where all important methods are explained.

In the mb_data repository, data that can be used for model fitting is contained. It has a file named fit_models.py that contains specifications of data and corresponding models. It can be called with one of these specifications to create a mdl file. In the directory data_models all mmdl files are currently stored that can be displayed by the lumen frontend.

The first step would be to implement a model there that is built by using a PPL. Then make sure that all necessary methods are correctly implemented and the visualization works as expected. One of the advantages of PPLs is the flexibility of the models specified, which is nullified by describing one fixed model. Therefore, the next steep is about thinking about a possibility to pass more flexible models. Maybe a function where you just give a model object as input that then automatically generates the necessary methods for the use in lumen?? Other possibilities? (would be good to ask that at the presentation)

compare to the plotting methods of PyMC3 / Stan

5.2 Requirements for a PPL

possible criterium: variety of distributions that can be described?

6 Konzept/Planung

In Bayesian data analysis we have a prior distribution for the parameters, $p(\theta)$, which represents our knowledge about them without seeing any data. We have also the likelihood of our data, $p(X|\theta)$, which represents the probability of obtaining the data X given the parameter θ. Using the Bayes' rule, we can then compute the posterior distribution for θ, $p(X|\theta)$:

$$p(\theta|X) = p(X|\theta) * p(\theta)/p(X) \tag{12}$$

Since the posterior distribution for θ can often not be computed analytically, Probabilistic Programming Languages use sampling methods to approximate the posterior distribution for θ.

When we want to visualize parameters and data, and also want to condition and marginalize arbitrarily, we need the joint distribution $p(x,\theta|X)$. The joint distribution can be computed by:

$$p(x,\theta|X) = p(x|\theta) * p(\theta|X) \tag{13}$$

The posterior distribution of the parameters, $p(\theta|X)$, can be approximated by PPLs through sampling, as stated above. (−> Make sure that this is true!!). To get density values from samples, we have to perform density estimation methods `http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KernelDensity.html#sklearn.neighbors.KernelDensity.score_samples`

7 Practical implementation

In this section, the choices and attempts made during the practical implementation of the interface are elucidated. At first, the choice of the PPL is justified.

My first approach will be to write a fixed probabilistic models as a model class, that fits one specific model to a specific set of data. (linear regression)

simple, fixed example: In Lumen, several methods have to be specified for a model class to work. The first, which will always be called first, is the method that assigns data to the model.

The second is the method that fits the model to the data. It will always be called second, after the data setting method. In Bayesian statistics, fitting means finding the joint posterior distribution. In the context of PPL, it means drawing samples from all variables, so that the joint posterior distribution can be approximated. The whole model should probably be specified in this method? Or before?

Another method is the one that marginalizes a random variable out. In a Bayesian model, a random variable can be a parameter as well as a data dimension. What is wanted here is thus the posterior distribution over all variables listed in keep. In PyMC3, marginalizing means sampling over a subset of the variables. Samples were already drawn in the fitting method, so only the corresponding

columns of those samples have to be chosen here.

Another method is the one that returns the probability density for a given point.

conditioning: posterior distribution given the values for the variables in remove. (Bad) Idea in PyMC3: Set the parameters in remove to a fixed value during the model specification, then sample as normal.

First Bug in the model run: Beim Laufen der Skripte für Lumen wir die Marginalisierungsmethode mit den Datendimensionen als Parameter angewendet. Meine erste Implementierung lässt aber nur Marginalisierungen auf Parameter zu, dadurch gibt es einen Fehler. Second Bug in the model: The marginalization method needs to access all variables of the data. However, it gets only the data passed for those variables to keep. So it cannot access the other variables. Problem: Model data is reduced to variables in keep. Lösungsidee: Setze zu Beginn der Methode die Daten auf den vollständigen Datensatz, und setze am Ende der Methode die Daten wieder auf die Dimensionen in keep zurück.

Problem: Es ist nicht gut, dass bei jedem Aufruf der _density-Methode gesampelt wird, da das viel Zeit und Speicher braucht. Lösung: Führe das Sampling nur einmal durch und speichere die samples in einer Klassenvariable.

7.1 Choose the PPL for the task at hand

8 Case examples

Here, case examples are presented to show how the finished interface performs in practice.

9 Conclusion

In this section, the result of the thesis is summarized (and maybe also the workflow leading to it?). It is possible to understand the main ideas and results of the thesis by reading only the introduction and this section.

Abbildungsverzeichnis

10 Literatur

Literatur

[1] J. Salvatier, T. V. Wiecki, and C. Fonnesbeck, "Probabilistic programming in python using PyMC3," *PeerJ Computer Science*, vol. 2, p. e55, apr 2016.

[2] A. Pfeffer, *Practical Probabilistic Programming*. Manning Publications, 2016.

[3] D. B. D. A. V. John B. Carlin, Hal S. Stern and D. B. R. A. Gelman, *Bayesian Data Analysis, 3Rd Edn*. T&F/Crc Press, 2014.

[4] C. Wang and D. M. Blei, "A general method for robust bayesian modeling," *Bayesian Analysis*, jan 2018.

[5] Wikipedia contributors, "Probabilistic programming language — Wikipedia, the free encyclopedia," 2018. [Online; accessed 23-August-2018].

[6] L. Hardesty, "Probabilistic programming does in 50 lines of code what used to take thousands," Apr. 2015. [Online; accessed 23-August-2018].

[7] "probabilistic-programming.org." [Online; accessed 23-August-2018].

[8] A. Gelman, D. Lee, and J. Guo, "Stan: A probabilistic programming language for bayesian inference and optimization," *Journal of Educational and Behavioral Statistics*, vol. 40, no. 5, pp. 530–543, 2015.

[9] S. Hoover, "presentations." `https://github.com/stephen-hoover/presentations`, 2016.